

Introduction to Machine Learning

Linear Regression

Nils M. Kriege

WS 2023

Data Mining and Machine Learning

Faculty of Computer Science

University of Vienna

Summary so far

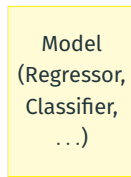

- Two basic forms of learning:
 - Supervised vs. Unsupervised learning
- Key challenge in ML:
 - Trading goodness of fit and model complexity
- Representation of data is of key importance

Supervised Learning Pipeline


Training data



$\mathcal{X} \quad \mathcal{Y}$



Test data



$f: \mathcal{X} \rightarrow \mathcal{Y}$

Representation

Model fitting

Prediction and
Generalization

- Instance of supervised learning
- **Goal:** Predict **real valued** labels (possibly vectors)
- Examples:


\mathcal{X}	\mathcal{Y}
Flight routes	delay (minutes)
Real estate objects	price
Patient & drug	treatment effectiveness
...	

More concrete real-world example: Diabetes

[Efron et al '04]


- Features \mathcal{X} :
 - Age
 - Sex
 - Body mass index
 - Average blood pressure
 - Six blood serum measurements (S1-S6)
- Label (target) \mathcal{Y} :
 - Quantitative measure of disease progression

Regression



- Goal: learn real valued mapping $f: \mathbb{R}^d \rightarrow \mathbb{R}$



Regression



- *Goal:* learn real valued mapping $f: \mathbb{R}^d \rightarrow \mathbb{R}$


Important choices in regression

- 1 What **types of functions f** should we consider? Examples:



- 2 How should we measure **goodness of fit**?

Example: Linear Regression



Goal: $y \approx f(x)$

Linear functions:

- 1-dim: $f(x) = ax + b$
- 2-dim: $f(x_1, x_2) = ax_1 + bx_2 + c$
- ...
- d -dim: $f(\mathbf{x}) = w_1x_1 + w_2x_2 + \cdots + w_dx_d + w_0 = \mathbf{w}^T\mathbf{x} + w_0$

Homogeneous Representation

Goal: Simplify $\mathbf{w}^T \mathbf{x} + w_0$ by removing the term w_0

Idea: Augment input data \mathbf{x} by adding a component $x_0 = 1$


$$\mathbf{w}^T \mathbf{x} + w_0 = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}$$

$$\mathbf{x} = (x_1, \dots, x_d)^T \quad \tilde{\mathbf{x}} = (1, x_1, \dots, x_d)^T$$

$$\mathbf{w} = (w_1, \dots, w_d)^T \quad \tilde{\mathbf{w}} = (w_0, w_1, \dots, w_d)^T$$

Quantifying goodness of fit

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\} \quad \mathbf{x}_i \in \mathbb{R}^d, y_i \in \mathbb{R}$$




Error at point i :

$$r_i = y_i - \mathbf{w}^T \mathbf{x}_i$$

Quantifying goodness of fit

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\} \quad \mathbf{x}_i \in \mathbb{R}^d, y_i \in \mathbb{R}$$



Error at point i :

$$r_i = y_i - \mathbf{w}^T \mathbf{x}_i$$

- Measure squared error per data point: $(y_i - \mathbf{w}^T \mathbf{x}_i)^2$
- Sum of errors

$$\hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Least-squares linear regression optimization


[Legendre 1805, Gauss 1809]

- Given a dataset $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$...
- ... how do we find the **optimal weight vector**?

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

How to solve? Example: Scikit Learn

```
# Create linear regression object  
regr = linear_model.LinearRegression()  
  
# Train the model using the training set  
regr.fit(X_train, Y_train)  
  
# Make predictions on the testing set  
Y_pred = regr.predict(X_test)
```



Method 1: Closed form solution

- Setting: $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ $\mathbf{x}_i \in \mathbb{R}^d, y_i \in \mathbb{R}$
- The problem

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

can be solved in **closed form**:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y},$$

where:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_n^T \end{pmatrix} = \begin{pmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,d} \\ x_{2,1} & x_{2,2} & \dots & x_{2,d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n,1} & x_{n,2} & \dots & x_{n,d} \end{pmatrix} \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}$$

Method 1: Closed form solution - Derivation

? How to derive the closed form solution?

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2 = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\begin{aligned} \hat{R}(\mathbf{w}) &= \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2 \\ &= \sum_{i=1}^n (\mathbf{y} - \mathbf{X}\mathbf{w})_i^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \\ &= \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\mathbf{w} + \mathbf{w}^T \mathbf{X}^T \mathbf{X}\mathbf{w} \end{aligned}$$

$$\nabla_{\mathbf{w}} \hat{R}(\mathbf{w}) = \mathbf{0} - 2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X}\mathbf{w}$$

$$\nabla_{\mathbf{w}} \hat{R}(\mathbf{w}) = \mathbf{0} \iff \mathbf{X}^T \mathbf{X}\mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\implies \hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Method 2: Optimization

- The objective function

$$\hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

is **convex**!

- A function $f: \mathbb{R}^d \rightarrow \mathbb{R}$ is **convex** iff $\forall \mathbf{x}, \mathbf{x}' \in \mathbb{R}^d, \lambda \in [0, 1]$, it holds that

$$f(\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}') \leq \lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{x}').$$

Gradient Descent

- Start at an arbitrary $\mathbf{w}_0 \in \mathbb{R}^d$
- For $t = 0, 1, 2, \dots$, do
 - $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla \hat{R}(\mathbf{w}_t)$
- Hereby, η_t is called **learning rate** or step rate

Convergence of gradient descent

- Under mild assumptions, if step size sufficiently small, gradient descent converges to a **stationary point** (gradient = 0)
- For convex objectives, it therefore finds the **optimal solution!**
- In the case of the squared loss, **linear convergence** for properly chosen **constant stepsize**.

Computing the gradient

$$\nabla \hat{R}(\mathbf{w}) = \left[\frac{\partial}{\partial w_1} \hat{R}(\mathbf{w}), \frac{\partial}{\partial w_2} \hat{R}(\mathbf{w}), \dots, \frac{\partial}{\partial w_d} \hat{R}(\mathbf{w}) \right]$$

$$\begin{aligned} \text{1-dim: } \nabla \hat{R}(w) &= \frac{\partial}{\partial w} \hat{R}(w) = \frac{\partial}{\partial w} \sum_{i=1}^n (y_i - wx_i)^2 \\ &= \sum_{i=1}^n \frac{\partial}{\partial w} (y_i - wx_i)^2 \quad \text{e.g., via chain rule} \\ &= \sum_{i=1}^n 2 \underbrace{(y_i - wx_i)}_{=r_i} (-x_i) = -2 \sum_{i=1}^n r_i x_i \end{aligned}$$

$$\text{d-dim: } \nabla \hat{R}(\mathbf{w}) = \dots = \sum_{i=1}^n 2(y_i - \mathbf{w}^T \mathbf{x}_i)(-\mathbf{x}_i) = -2 \sum_{i=1}^n r_i \mathbf{x}_i$$

Demo: Gradient descent

Choosing a stepsize

What happens if we choose a poor stepsize?

Adaptive step size

- Can update the step size adaptively. For example:
 - (a) Via **line search** (optimizing step size every step):

Suppose that at iteration t , we have $\mathbf{w}_t, \mathbf{g}_t = \nabla \hat{R}(\mathbf{w}_t)$

Define: $\eta_t^* = \arg \min_{\eta \in [0, \infty]} \hat{R}(\mathbf{w}_t - \eta \mathbf{g}_t)$

- (b) “Bold driver” heuristic:
 - Initial learning rate η_0 .
 - If function decreases, increase step size:

If $\hat{R}(\mathbf{w}_{t+1}) < \hat{R}(\mathbf{w}_t)$ then $\eta_{t+1} = \eta_t c_+$ with $c_+ > 1$.

- If function increases, decrease step size:

If $\hat{R}(\mathbf{w}_{t+1}) > \hat{R}(\mathbf{w}_t)$ then $\eta_{t+1} = \eta_t c_-$ with $c_- < 1$.

Demo: Gradient Descent for Linear Regression

Gradient Descent vs Closed Form

Why would one ever consider performing gradient descent, when it is possible to find closed form solution?

Closed form:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Gradient descent:

$$\nabla \hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

Gradient Descent vs Closed Form

Why would one ever consider performing gradient descent, when it is possible to find closed form solution?

Closed form:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Gradient descent:

$$\nabla \hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

- Computational complexity

Gradient Descent vs Closed Form

Why would one ever consider performing gradient descent, when it is possible to find closed form solution?

Closed form:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Gradient descent:

$$\nabla \hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

- Computational complexity
- May not need an optimal solution

Gradient Descent vs Closed Form

Why would one ever consider performing gradient descent, when it is possible to find closed form solution?

Closed form:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$


Gradient descent:

$$\nabla \hat{R}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

- Computational complexity
- May not need an optimal solution
- Many problems don't admit closed form solution

Other Loss Functions

- So far: Measure goodness of fit via squared error
- Many other **loss functions** possible (and sensible!)



Supervised learning summary so far

Representation/
features

Linear hypotheses

Model/
objective

Loss-function (squared loss, ℓ_p loss)


Method

Exact solution, Gradient Descent


Evaluation
metric

Empirical risk = (mean) squared error

Outlook: Fitting nonlinear functions



Outlook: Fitting nonlinear functions



Outlook: Linear regression for polynomials

We can fit non-linear functions via linear regression, using nonlinear features of our data (basis functions):

$$f(\mathbf{x}) = \sum_{i=1}^D w_i \phi_i(\mathbf{x})$$