

# Introduction to Machine Learning

## Generalization and Model Validation

---

Nils M. Kriege


WS 2023

Data Mining and Machine Learning

Faculty of Computer Science


University of Vienna

# Regression




- Goal: learn real valued mapping  $f: \mathbb{R}^d \rightarrow \mathbb{R}$

# Regression



- *Goal:* learn real valued mapping  $f: \mathbb{R}^d \rightarrow \mathbb{R}$


# Fitting nonlinear functions



Goal:  $y \approx f(x)$

Observation: Linear  
function not suitable

# Fitting nonlinear functions




Goal:  $y \approx f(x)$

Observation: Linear  
function not suitable



Use quadratic function  $f(x) = ax^2 + bx + c$

# Fitting nonlinear functions



Goal:  $y \approx f(x)$

Observation: Linear  
function not suitable



Use quadratic function  $f(x) = ax^2 + bx + c$



Transform input data and use linear regression:

$$f(x) = ax^2 + bx + c = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}} \quad \text{with} \quad \tilde{\mathbf{w}} = \begin{pmatrix} a \\ b \\ c \end{pmatrix}, \quad \tilde{\mathbf{x}} = \begin{pmatrix} x^2 \\ x \\ 1 \end{pmatrix}$$

# Linear regression for polynomials

We can fit **non-linear functions** via **linear regression**, using non-linear features of our data (**basis functions**):

$$\phi: \mathbb{R}^d \rightarrow \mathbb{R}^D \quad f(\mathbf{x}) = \sum_{i=1}^D w_i \phi_i(\mathbf{x})$$

**Example:** Polynomials of degree  $k$ .

$$\text{1-dim: } \phi(x) = (1, x, x^2, \dots, x^k)^T$$





$$\text{2-dim: } \phi(\mathbf{x}) = (1, x_1, x_2, x_1x_2, x_1^2, x_2^2, x_1^2x_2, \dots)^T$$

$$\text{\textit{d}-dim: } \phi(\mathbf{x}) = \text{vector of all monomials } x_1, x_2, \dots, x_d \text{ of degree } k$$


# Demo: Linear regression on polynomials





# Observation: Underfitting and overfitting




# Observation: Underfitting and overfitting



⚠ Underfitting



⚠ Overfitting




⚠ Overfitting

## Supervised learning summary so far

Representation/ features	Linear hypotheses, nonlinear hypotheses through feature transformations
Model/ objective	Loss-function (squared loss, $\ell_p$ loss)
Method	Exact solution, Gradient Descent
Evaluation metric	Empirical risk = (mean) squared error


## Model selection for linear regression with polynomials



## Interlude: A note on probability

- You'll need to know about basic concepts in probability:
  - Random variables
  - Expectations (Mean, variance etc.)
  - Independence (i.i.d. samples from a distribution, ...)
  - ...


## Example: Gaussian distribution




- Standard deviation  $\sigma$
- Mean  $\mu$
- Probability density function  $p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$

# Example: Multivariate Gaussian

$$\frac{1}{2\pi\sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}(x-\mu)^T\Sigma^{-1}(x-\mu)\right) \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{21} & \sigma_2^2 \end{pmatrix} \quad \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}$$



$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$



$$\Sigma = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

:

## Interlude: Expectations

- Expected value of random variable  $X$  with probability density function  $p$ :

$$\mathbb{E}[X] = \int_{\mathcal{X}} xp(x) dx$$

- Expected value of some function of  $X$ :

$$\mathbb{E}[f(x)] = \int_{\mathcal{X}} f(x)p(x) dx$$

- Linearity of expectation:

$$\mathbb{E}[X + Y] = \mathbb{E}(X) + \mathbb{E}(Y)$$

$$\mathbb{E}[aX] = a\mathbb{E}[X]$$



# Achieving generalization

- **Fundamental assumption:** Our data set is generated **independently and identically distributed (iid)** from some **unknown** distribution  $P$ :

$$(\mathbf{x}_i, y_i) \sim P(\mathbf{X}, Y)$$

- Our goal is to minimize **the expected error (true risk)** under  $P$ :

$$\begin{aligned} R(\mathbf{w}) &= \int P(\mathbf{x}, y)(y - \mathbf{w}^T \mathbf{x})^2 d\mathbf{x} dy \\ &= \mathbb{E}_{\mathbf{x}, y}[(y - \mathbf{w}^T \mathbf{x})^2] \end{aligned}$$

## Side note on iid assumption

- When is iid assumption invalid?
  - Time series data
  - Spatially correlated data
  - Correlated noise
  - ...
- Often, can still use machine learning, but one has to be careful in interpreting results.
- Most important: Choose train/test to assess the desired generalization

# Estimating the generalization error

- Estimate the **true risk** by the **empirical risk** on a sample data set  $\mathcal{D}$ :

$$\hat{R}_{\mathcal{D}}(\mathbf{w}) = \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, y) \in \mathcal{D}} (y - \mathbf{w}^T \mathbf{x})^2$$

- Why might this work?

# Estimating the generalization error

- Estimate the **true risk** by the **empirical risk** on a sample data set  $\mathcal{D}$ :

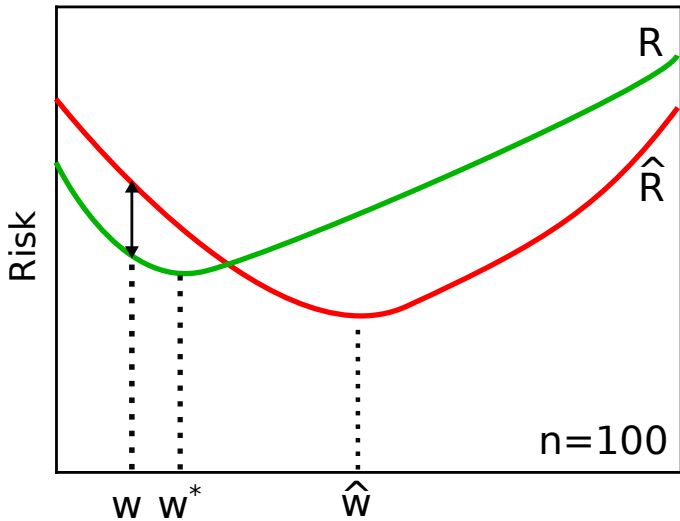
$$\hat{R}_{\mathcal{D}}(\mathbf{w}) = \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, y) \in \mathcal{D}} (y - \mathbf{w}^T \mathbf{x})^2$$

- Why might this work?  
**Law of large numbers:**  $\hat{R}_{\mathcal{D}}(\mathbf{w}) \rightarrow R(\mathbf{w})$  for any fixed  $\mathbf{w}$  almost surely as  $|\mathcal{D}| \rightarrow \infty$


# What happens if we optimize on training data?

- Suppose we are given training data  $\mathcal{D}$
- **Empirical Risk Minimization:**  $\hat{\mathbf{w}}_{\mathcal{D}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})$
- Ideally we wish to solve:  $\mathbf{w}^* = \arg \min_{\mathbf{w}} R(\mathbf{w})$

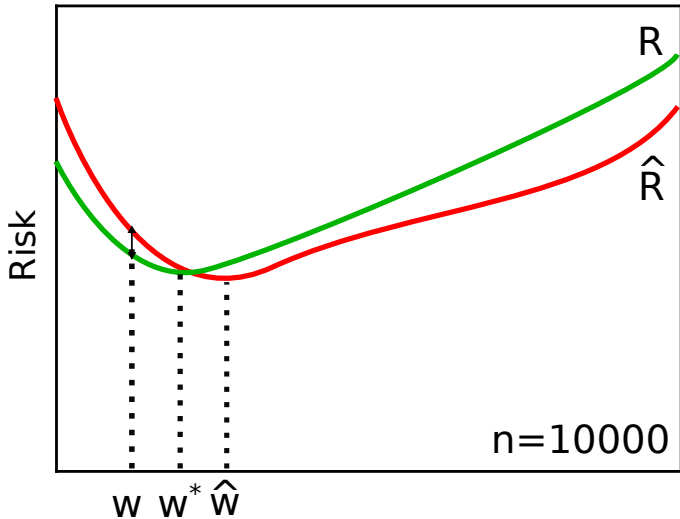
## Empirical vs true risk



## Empirical vs true risk



## Empirical vs true risk





## Outlook: Requirements for learning


- For learning via empirical risk minimization to be successful, need **uniform convergence**:

$$\sup_{\mathbf{w}} |R(\mathbf{w}) - \hat{R}_{\mathcal{D}}(\mathbf{w})| \rightarrow 0 \quad \text{as} \quad |\mathcal{D}| \rightarrow \infty$$

- This is **not** implied by law of large numbers alone, but depends on model class (holds, e.g., for squared loss on data distributions with bounded support)  
→ Statistical learning theory

# What can go wrong? An example

Consider the following classification problem:




- $\mathbf{x}_i \in \mathbb{R}^2$ ,  $y_i \in \{\text{o}, \text{x}\}$
- Points in  $\square$  are  $\text{o}$  and  $\text{x}$  otherwise.
- $\text{area}(\square) = 2 \cdot \text{area}(\square)$
- A point drawn uniformly at random is  $\text{o}$  or  $\text{x}$  with the same probability.

- Training data:  $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- Classifier:  $h'_{\mathcal{D}}(\mathbf{x}) = \begin{cases} y & \text{if } (\mathbf{x}, y) \in \mathcal{D}, \\ \text{o} & \text{otherwise.} \end{cases}$
- Empirical risk of  $h'_{\mathcal{D}}$ :      True risk of  $h'_{\mathcal{D}}$ :

# What can go wrong? An example


Consider the following classification problem:




- $\mathbf{x}_i \in \mathbb{R}^2$ ,  $y_i \in \{\text{o}, \text{x}\}$
- Points in  $\square$  are  $\text{o}$  and  $\text{x}$  otherwise.
- $\text{area}(\square) = 2 \cdot \text{area}(\square)$
- A point drawn uniformly at random is  $\text{o}$  or  $\text{x}$  with the same probability.

- Training data:  $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- Classifier:  $h'_{\mathcal{D}}(\mathbf{x}) = \begin{cases} y & \text{if } (\mathbf{x}, y) \in \mathcal{D}, \\ \text{o} & \text{otherwise.} \end{cases}$
- Empirical risk of  $h'_{\mathcal{D}}$ : 0      True risk of  $h'_{\mathcal{D}}$ :  $\frac{1}{2}$


## What can go wrong in ERM?



## What can go wrong in ERM?



## What can go wrong in ERM?



- Law of large numbers / uniform convergence are **asymptotic** statements (with  $n \rightarrow \infty$ )
- In practice one has **finite** amount of data.
- What can go wrong?

## Simple example

$$\hat{\mathbf{w}}_{\mathcal{D}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} R(\mathbf{w})$$



## What if we evaluate performance on training data?

$$\hat{\mathbf{w}}_{\mathcal{D}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} R(\mathbf{w})$$

- In general it holds that  $\mathbb{E}_{\mathcal{D}}[\hat{R}_{\mathcal{D}}(\hat{\mathbf{w}}_{\mathcal{D}})] \leq \mathbb{E}_{\mathcal{D}}[R(\hat{\mathbf{w}}_{\mathcal{D}})]$

## What if we evaluate performance on training data?

$$\hat{\mathbf{w}}_{\mathcal{D}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} R(\mathbf{w})$$

- In general it holds that  $\mathbb{E}_{\mathcal{D}}[\hat{R}_{\mathcal{D}}(\hat{\mathbf{w}}_{\mathcal{D}})] \leq \mathbb{E}_{\mathcal{D}}[R(\hat{\mathbf{w}}_{\mathcal{D}})]$
- Thus we obtain an overly optimistic estimate!

$$\hat{\mathbf{w}}_{\mathcal{D}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} R(\mathbf{w})$$

$$\mathbb{E}_{\mathcal{D}}[\hat{R}_{\mathcal{D}}(\hat{\mathbf{w}}_{\mathcal{D}})] = \mathbb{E}_{\mathcal{D}}[\min_{\mathbf{w}} \hat{R}_{\mathcal{D}}(\mathbf{w})] \quad (\text{ERM})$$

$$\leq \min_{\mathbf{w}} \mathbb{E}_{\mathcal{D}}[\hat{R}_{\mathcal{D}}(\mathbf{w})] \quad (\text{Jensen's inequality})$$

$$= \min_{\mathbf{w}} \mathbb{E}_{\mathcal{D}} \left[ \frac{1}{|\mathcal{D}|} \sum_{i=1}^{\mathcal{D}} (y_i - \mathbf{w}\mathbf{x}_i)^2 \right] \quad (\text{Def. } \hat{R}_{\mathcal{D}}(\mathbf{w}))$$

$$= \min_{\mathbf{w}} \frac{1}{|\mathcal{D}|} \sum_{i=1}^{\mathcal{D}} \underbrace{\mathbb{E}_{(\mathbf{x}_i, y_i) \sim P} [(y_i - \mathbf{w}\mathbf{x}_i)^2]}_{R(\mathbf{w})} \quad (\text{lin. exp.})$$

$$= \min_{\mathbf{w}} R(\mathbf{w}) \leq \mathbb{E}_{\mathcal{D}}[R(\hat{\mathbf{w}}_{\mathcal{D}})]$$

## More realistic evaluation?

- Want to avoid underestimating the prediction error
- Idea: Use **separate** test set from the same distribution  $P$
- Obtain indep. training and test data  $\mathcal{D}_{\text{train}}$  and  $\mathcal{D}_{\text{test}}$
- Optimize  $\mathbf{w}$  on training set:

$$\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}} = \arg \min_{\mathbf{w}} \hat{R}_{\mathcal{D}_{\text{train}}}(\mathbf{w})$$

- Evaluate on test set:

$$\hat{R}_{\mathcal{D}_{\text{test}}}(\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}}) = \frac{1}{|\mathcal{D}_{\text{test}}|} \sum_{(\mathbf{x}, y) \in \mathcal{D}_{\text{test}}} (y - \hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}}^T \mathbf{x})^2$$

- Then:

$$\mathbb{E}_{\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{test}}} [\hat{R}_{\mathcal{D}_{\text{test}}}(\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}})] = \mathbb{E}_{\mathcal{D}_{\text{train}}} [R(\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}})]$$

# Why?

---

## Recap: Evaluating predictive performance

- Training error (empirical risk) **systematically underestimates** true risk:

$$\mathbb{E}_{\mathcal{D}}[\hat{R}_{\mathcal{D}}(\hat{\mathbf{w}}_{\mathcal{D}})] \leq \mathbb{E}_{\mathcal{D}}[R(\hat{\mathbf{w}}_{\mathcal{D}})]$$

- Using an **independent test set** avoids this bias:

$$\mathbb{E}_{\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{test}}}[\hat{R}_{\mathcal{D}_{\text{test}}}(\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}})] = \mathbb{E}_{\mathcal{D}_{\text{train}}}[R(\hat{\mathbf{w}}_{\mathcal{D}_{\text{train}}})]$$

## First Attempt: Evaluation for Model Selection

- Obtain training and test data  $\mathcal{D}_{\text{train}}$  and  $\mathcal{D}_{\text{test}}$
- Fit each candidate model (e.g., degree  $m$  of polynomial):

$$\hat{\mathbf{w}}_m = \underset{\mathbf{w}: \text{degree}(\mathbf{w}) \leq m}{\text{argmin}} \hat{R}_{\text{train}}(\mathbf{w})$$

- Pick one which does best on test set:

$$\hat{m} = \underset{m}{\text{argmin}} \hat{R}_{\text{test}}(\hat{\mathbf{w}}_m)$$

## First Attempt: Evaluation for Model Selection

- Obtain training and test data  $\mathcal{D}_{\text{train}}$  and  $\mathcal{D}_{\text{test}}$
- Fit each candidate model (e.g., degree  $m$  of polynomial):

$$\hat{\mathbf{w}}_m = \underset{\mathbf{w}: \text{degree}(\mathbf{w}) \leq m}{\text{argmin}} \hat{R}_{\text{train}}(\mathbf{w})$$


- Pick one which does best on test set:

$$\hat{m} = \underset{m}{\text{argmin}} \hat{R}_{\text{test}}(\hat{\mathbf{w}}_m)$$

- Do you see a problem?



## Overfitting to test set



- Test error is itself random! Variance usually increases for more complex models
- Optimizing for **single** test set creates bias

## Solution: Pick Multiple Test Sets!

- **Key idea:** Instead of using a single test set, use **multiple test sets** and average to decrease variance
- **Dilemma:** Any data I use for testing I can't use for training

## Solution: Pick Multiple Test Sets!

- **Key idea:** Instead of using a single test set, use **multiple test sets** and average to decrease variance
- **Dilemma:** Any data I use for testing I can't use for training
- $\Rightarrow$  **Using multiple independent test sets is expensive and wasteful**