

Introduction to Machine Learning

Non-linear Prediction with Kernels

Nils M. Kriege


WS 2023

Data Mining and Machine Learning

Faculty of Computer Science

University of Vienna

Recall: Linear classifiers



$$\hat{y} = \text{sign}(\mathbf{w}^T \mathbf{x})$$

Recall: The Perceptron problem

- Solve

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n \ell_P(\mathbf{w}; y_i, \mathbf{x}_i)$$

where

$$\ell_P(\mathbf{w}; y_i, \mathbf{x}_i) = \max(0, -y_i \mathbf{w}^T \mathbf{x}_i)$$

- Optimize via **Stochastic Gradient Descent**

Solving non-linear classification tasks

- How can we find nonlinear classification boundaries?


Solving non-linear classification tasks

- How can we find nonlinear classification boundaries?
- Similar as in regression, can use **non-linear transformations** of the feature vectors, followed by **linear classification**



$$x \mapsto [x_1, x_2]$$

\Rightarrow



Recall: Linear regression for polynomials

- We can fit **non-linear functions via linear regression**, using **nonlinear features** of our data (basis functions):

$$f(\mathbf{x}) = \sum_{i=1}^d w_i \phi_i(\mathbf{x})$$

- For example: polynomials (in 1d):

$$f(x) = \sum_{i=0}^m w_i x^i$$

Polynomials in higher dimensions


- Suppose we wish to use polynomial features, but our input is higher-dimensional
- Can still use monomial features
- **Example:** Monomials in 2 variables, degree = 2;

$$\mathbf{x} = [x_1, x_2] \quad \mapsto \quad \phi(\mathbf{x}) = [x_1^2, x_2^2, x_1x_2]$$



Avoiding the feature explosion

- Need $O(d^k)$ dimensions to represent (multivariate) polynomials of degree k on d features


Avoiding the feature explosion

- Need $O(d^k)$ dimensions to represent (multivariate) polynomials of degree k on d features
-  Example: $d = 10000$, $k = 2 \Rightarrow$ Need $\sim 100M$ dimensions

Avoiding the feature explosion


- Need $O(d^k)$ dimensions to represent (multivariate) polynomials of degree k on d features
-  Example: $d = 10000$, $k = 2 \Rightarrow$ Need $\sim 100M$ dimensions
-  In the following, we can see how we can efficiently **implicitly** operate in such high-dimensional feature spaces (i.e., without ever explicitly computing the transformation)

Revisiting the Perceptron/SVM

-  Fundamental insight: Optimal hyperplane **lies in the span of the data**:

$$\hat{\mathbf{w}} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$


Revisiting the Perceptron/SVM

-  Fundamental insight: Optimal hyperplane **lies in the span of the data**:

$$\hat{\mathbf{w}} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

- **(Handwavy) proof:** (Stochastic) gradient descent starting from 0 **constructs** such a representation:
 - Perceptron: $\mathbf{w}_{t+1} = \mathbf{w}_t + \eta_t y_t \mathbf{x}_t [y_t \mathbf{w}_t \mathbf{x}_t < 0]$
 - SVM: $\mathbf{w}_{t+1} = \mathbf{w}_t (1 - 2\lambda \eta_t) + \eta_t y_t \mathbf{x}_t [y_t \mathbf{w}_t \mathbf{x}_t < 1]$

Revisiting the Perceptron/SVM

-  Fundamental insight: Optimal hyperplane lies in the span of the data:

$$\hat{\mathbf{w}} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

- **(Handwavy) proof:** (Stochastic) gradient descent starting from 0 constructs such a representation:
 - Perceptron: $\mathbf{w}_{t+1} = \mathbf{w}_t + \eta_t y_t \mathbf{x}_t [y_t \mathbf{w}_t \mathbf{x}_t < 0]$
 - SVM: $\mathbf{w}_{t+1} = \mathbf{w}_t (1 - 2\lambda \eta_t) + \eta_t y_t \mathbf{x}_t [y_t \mathbf{w}_t \mathbf{x}_t < 1]$
- **More abstract proof:** Follows from so called “representer theorems” (discussed later)

Reformulating the Perceptron

Idea: Replace \mathbf{w} by $\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max(0, -y_i \mathbf{w}^T \mathbf{x}_i)$$

$$\hat{\alpha} = \underset{\alpha}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max \left\{ 0, -y_i \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j \right)^T \mathbf{x}_i \right\}$$

$$= \underset{\alpha}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max \left\{ 0, -y_i \sum_{j=1}^n \alpha_j y_j \mathbf{x}_j^T \mathbf{x}_i \right\}$$


$$= \underset{\alpha}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max \left\{ 0, - \sum_{j=1}^n \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \right\}$$

Advantage of reformulation

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j\}$$

Advantage of reformulation

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j\}$$

-  **Key observation:** Objective only depends on **inner products** of pairs of data points
- Thus, we can **implicitly** work in high-dimensional spaces, as long as we can do inner products efficiently:

$$\begin{aligned}\mathbf{x} &\mapsto \phi(\mathbf{x}) \\ \mathbf{x}^T \mathbf{x}' &\mapsto \phi(\mathbf{x})^T \phi(\mathbf{x}') =: k(\mathbf{x}, \mathbf{x}')\end{aligned}$$

“Kernels = efficient inner products”

- Often, $k(\mathbf{x}, \mathbf{x}')$ can be computed much more efficiently than $\phi(\mathbf{x})^T \phi(\mathbf{x})'$

“Kernels = efficient inner products”

- Often, $k(\mathbf{x}, \mathbf{x}')$ can be computed much more efficiently than $\phi(\mathbf{x})^T \phi(\mathbf{x})'$
- Simple example: Polynomial kernel in degree 2

$$\begin{aligned}\mathbf{x} &= [x_1, x_2]^T \mapsto \phi(\mathbf{x}) = [x_1^2, x_2^2, \sqrt{2}x_1x_2] \\ \phi(\mathbf{x})^T \phi(\mathbf{x}') &= x_1^2 x_1'^2 + x_2^2 x_2'^2 + 2x_1x_2x_1'x_2' \\ &= (\mathbf{x}^T \mathbf{x}')^2 =: k(\mathbf{x}, \mathbf{x}')\end{aligned}$$

Embedding: #+: 2 #*: 3+3+4=10

Functional: #+: 1 #*: 3

Polynomial kernels (degree 2)

- Suppose $\mathbf{x} = [x_1, \dots, x_d]^T$ and $\mathbf{x}' = [x'_1, \dots, x'_d]^T$
- Then

$$\begin{aligned}(\mathbf{x}^T \mathbf{x}')^2 &= \left(\sum_{i=1}^d x_i x'_i \right)^2 = \sum_{i=1}^d x_i^2 x'^2_i + 2 \sum_{1 \leq i < j \leq d} x_i x'_i x_j x'_j \\ &= \phi(\mathbf{x})^T \phi(\mathbf{x}'), \text{ where} \\ \phi(\mathbf{x}) &= [x_1^2, \dots, x_d^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \dots, \sqrt{2}x_{d-1}x_d]^T\end{aligned}$$

Polynomial kernels: Fixed degree

- The kernel $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of degree m

Polynomial kernels: Fixed degree

- The kernel $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of degree m

$$x_1^m, x_2^m, \dots, x_d^m, x_1^{m-1}x_2, \dots, x_1^{m-1}x_d, \dots, x_1 \cdots x_m, \dots, x_{d-m+1} \cdots x_d$$

Monomials of degree m in d variables:

$$\binom{d+m-1}{m} = O(d^m)$$

Polynomial kernels: Fixed degree

- The kernel $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of degree m

$$x_1^m, x_2^m, \dots, x_d^m, x_1^{m-1}x_2, \dots, x_1^{m-1}x_d, \dots, x_1 \cdots x_m, \dots, x_{d-m+1} \cdots x_d$$

Monomials of degree m in d variables:

$$\binom{d+m-1}{m} = O(d^m)$$

- How can we get monomials **up to** order m ?

- The polynomial kernel $k(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of up to degree m

- The polynomial kernel $k(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of up to degree m

$$1, x_1, x_2, \dots, x_d, x_1^2, x_2^2, \dots, x_d^2, x_1 x_2 \cdots x_m, \dots$$

Monomials of degree m in d variables:

$$\binom{d+m}{m} = O(d^m)$$

Polynomial kernels

- The polynomial kernel $k(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x}^T \mathbf{x}')^m$ implicitly represents all monomials of up to degree m

$$1, x_1, x_2, \dots, x_d, x_1^2, x_2^2, \dots, x_d^2, x_1 x_2 \cdots x_m, \dots$$

Monomials of degree m in d variables:

$$\binom{d+m}{m} = O(d^m)$$

-  Representing the monomials (and computing inner product explicitly) is exponential in m !

The “Kernel Trick”

Kernel Trick

- Express problem s.t. it only depends on inner products
- Replace inner products by kernels

The “Kernel Trick”

Kernel Trick

- Express problem s.t. it only depends on inner products
- Replace inner products by kernels

$$\mathbf{x}_i^T \mathbf{x}_j \Rightarrow k(\mathbf{x}_i, \mathbf{x}_j)$$

- This “trick” is very widely applicable!

The “Kernel Trick”

Kernel Trick

- Express problem s.t. it only depends on inner products
 - Replace inner products by kernels
-
- Example: Perceptron

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j\}$$



$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)\}$$

The “Kernel Trick”

Kernel Trick

- Express problem s.t. it only depends on inner products
 - Replace inner products by kernels
-
- Example: Perceptron

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j\}$$



$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -\sum_{j=1}^n \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)\}$$

- Will see further examples later

Derivation: Kernelized Perceptron

Perceptron

Training:

$$\mathbf{w}_0 \leftarrow \mathbf{0}$$

For $t = 0, 1, 2, \dots$ do

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t$$

Sample $(\mathbf{x}_i, y_i) \in \mathcal{D}$

if $y_i \mathbf{w}_t^T \mathbf{x}_i < 0$

$$\mathbf{w}_{t+1} \leftarrow \eta_t \mathbf{w}_{t+1} + y_i \mathbf{x}_i$$

Prediction:

$$\text{sign}(\mathbf{w}^T \mathbf{x})$$

Derivation: Kernelized Perceptron

Perceptron

Training:

$$\mathbf{w}_0 \leftarrow \mathbf{0}$$

For $t = 0, 1, 2, \dots$ do

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t$$

Sample $(\mathbf{x}_i, y_i) \in \mathcal{D}$

if $y_i \mathbf{w}_t^T \mathbf{x}_i < 0$

$$\mathbf{w}_{t+1} \leftarrow \eta_t \mathbf{w}_{t+1} + y_i \mathbf{x}_i$$

Prediction:

$$\text{sign}(\mathbf{w}^T \mathbf{x})$$

Kernelized Perceptron

Training:

$$\alpha^{(0)} \leftarrow \mathbf{0}$$

For $t = 0, 1, 2, \dots$ do

$$\alpha^{(t+1)} \leftarrow \alpha^{(t)}$$

Sample $(\mathbf{x}_i, y_i) \in \mathcal{D}$

if $y_i \sum_{j=1}^n \alpha_j^{(t)} y_j \mathbf{x}_j^T \mathbf{x}_i < 0$

$$\alpha_i^{(t+1)} \leftarrow \alpha_i^{(t+1)} + \eta_t$$

Prediction:

$$\text{sign}(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j^T \mathbf{x})$$

Kernelized Perceptron

Training

- Initialize $\alpha_1 = \dots = \alpha_n = 0$
- For $t = 1, 2, \dots$
 - Pick data point (\mathbf{x}_i, y_i) uniformly at random
 - Predict

$$\hat{y} = \text{sign}\left(\sum_{j=1}^n \alpha_j y_j k(\mathbf{x}_j, \mathbf{x}_i)\right)$$

- if $\hat{y} \neq y_i$ set $\alpha_i \leftarrow \alpha_i + \eta_t$

Prediction

- For new point \mathbf{x} predict

$$\hat{y} = \text{sign}\left(\sum_{j=1}^n \alpha_j y_j k(\mathbf{x}_j, \mathbf{x})\right)$$

Demo: Kernelized Perceptron

The Kernel Trick: Summary

Regression/Classification
problem in \mathbf{x}, y
which cannot be
separated well linearly

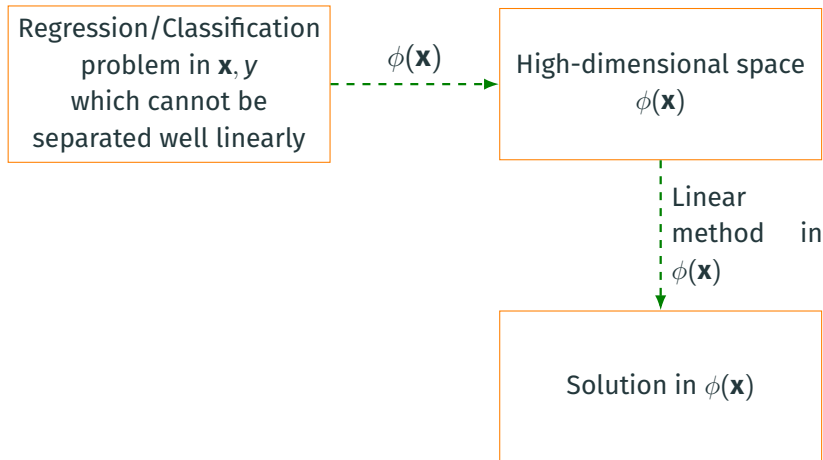
The Kernel Trick: Summary

Regression/Classification
problem in \mathbf{x}, y
which cannot be
separated well linearly

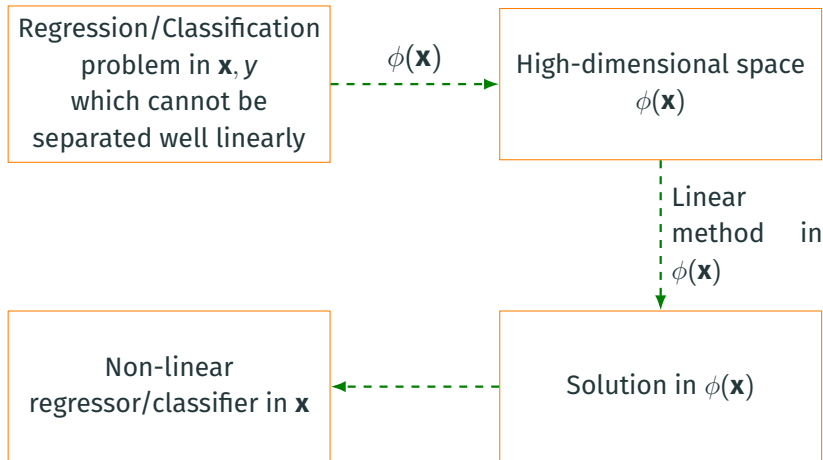
$\phi(\mathbf{x})$

High-dimensional space
 $\phi(\mathbf{x})$

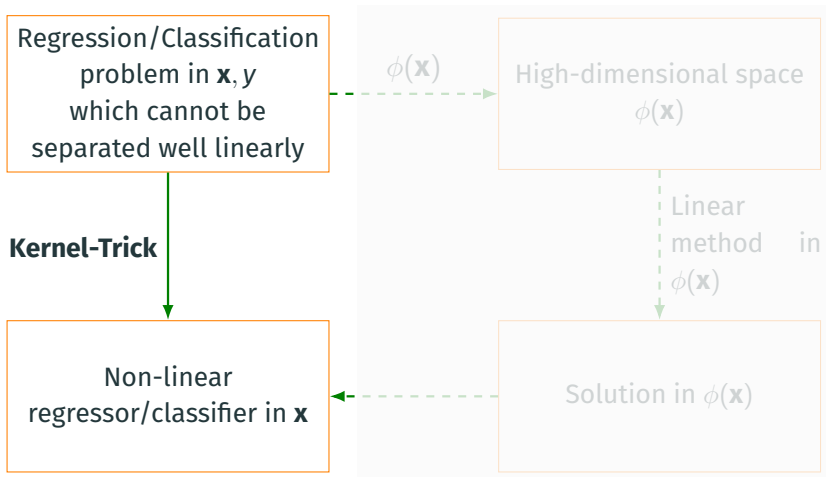
The Kernel Trick: Summary



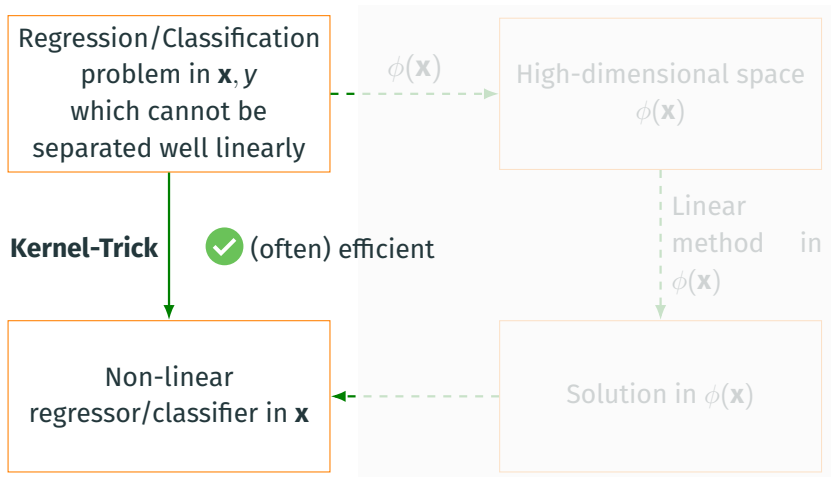
The Kernel Trick: Summary



The Kernel Trick: Summary



The Kernel Trick: Summary



- What are valid kernels?
- How can we select a good kernel for our problem?
- Can we use kernels beyond the perceptron?
- Kernels work in very high-dimensional spaces. Doesn't this lead to overfitting?

Properties of kernel functions

- Data space \mathcal{X}
- A kernel is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ? Can we use any function?

Properties of kernel functions

- Data space \mathcal{X}
- A kernel is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ? Can we use any function?
- k must be an inner product in a suitable space

Properties of kernel functions

- Data space \mathcal{X}
- A kernel is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ? Can we use any function?
- k must be an **inner product** in a suitable space
 $\Rightarrow k$ must be **symmetric**!

$$\forall \mathbf{x}, \mathbf{x}' \in \mathcal{X}: k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}') = \phi(\mathbf{x}')^T \phi(\mathbf{x}) = k(\mathbf{x}', \mathbf{x})$$

Properties of kernel functions

- Data space \mathcal{X}
- A kernel is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ? Can we use any function?
- k must be an **inner product** in a suitable space
 $\Rightarrow k$ must be **symmetric**!

$$\forall \mathbf{x}, \mathbf{x}' \in \mathcal{X}: k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}') = \phi(\mathbf{x}')^T \phi(\mathbf{x}) = k(\mathbf{x}', \mathbf{x})$$

\Rightarrow Are there any other properties that it must satisfy?

Positive semi-definite matrices

Symmetric matrix $M \in \mathbb{R}^{n \times n}$ is **positive semidefinite** iff any of the following two conditions holds

1. $\forall \mathbf{x} \in \mathbb{R}^n : \mathbf{x}^T M \mathbf{x} \geq 0$
2. all eigenvalues of M are ≥ 0

Kernels \Rightarrow Semi-definite matrices

- Data space \mathcal{X} (possibly finite)
- Kernel $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ with feature map $\phi: \mathcal{X} \rightarrow \mathbb{R}^d$
- Take any finite subset of data $S = \{\mathbf{x}_1, \dots, \mathbf{x}_n\} \subseteq \mathcal{X}$
- Then the **kernel (gram) matrix**

$$\mathbf{K} = \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix} = \begin{pmatrix} \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_1) & \cdots & \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_n) \\ \vdots & & \vdots \\ \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_1) & \cdots & \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_n) \end{pmatrix}$$

is **positive semidefinite**

Kernels \Rightarrow Semi-definite matrices

- Data space \mathcal{X} (possibly finite)
- Kernel $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ with feature map $\phi: \mathcal{X} \rightarrow \mathbb{R}^d$
- Take any finite subset of data $S = \{\mathbf{x}_1, \dots, \mathbf{x}_n\} \subseteq \mathcal{X}$
- Then the **kernel (gram) matrix**

$$\mathbf{K} = \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix} = \begin{pmatrix} \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_1) & \cdots & \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_n) \\ \vdots & & \vdots \\ \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_1) & \cdots & \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_n) \end{pmatrix}$$

is **positive semidefinite**

Proof.

$\mathbf{K} = \Phi^T \Phi$, where $\Phi = [\phi(\mathbf{x}_1), \phi(\mathbf{x}_2), \dots, \phi(\mathbf{x}_n)]$ ($d \times n$ matrix)

$$\mathbf{x}^T \mathbf{K} \mathbf{x} = \mathbf{x}^T (\Phi^T \Phi) \mathbf{x} = (\mathbf{x}^T \Phi^T) (\Phi \mathbf{x}) = (\Phi \mathbf{x})^T (\Phi \mathbf{x}) = \|\Phi \mathbf{x}\|_2^2 \geq 0 \quad \square$$

Semi-definite matrices \Rightarrow Kernels

- Suppose the data space $\mathcal{X} = \{1, \dots, n\}$ is **finite**, and we are given a positive semidefinite matrix $\mathbf{K} \in \mathbb{R}^{n \times n}$
- Then we **can always construct a feature map**

$$\phi: \mathcal{X} \rightarrow \mathbb{R}^n$$

$$\text{such that } \mathbf{K}_{i,j} = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

Proof.

Since \mathbf{K} is psd, its eigendecomposition $\mathbf{K} = \mathbf{U}\mathbf{D}\mathbf{U}^T$ exists with $\mathbf{D} = \text{diag}(\lambda_1, \dots, \lambda_n)$ and $\lambda_i \geq 0$. Hence, $\mathbf{D} = \mathbf{D}^{1/2}(\mathbf{D}^{1/2})^T$ with $\mathbf{D}^{1/2} = \text{diag}(\sqrt{\lambda_1}, \dots, \sqrt{\lambda_n})$ for all $1 \leq i \leq n$.

$$\mathbf{K} = \mathbf{U}\mathbf{D}\mathbf{U}^T = \mathbf{U}(\mathbf{D}^{1/2}(\mathbf{D}^{1/2})^T)\mathbf{U}^T = \underbrace{(\mathbf{U}\mathbf{D}^{1/2})}_{=:\Phi^T} \underbrace{((\mathbf{D}^{1/2})^T \mathbf{U}^T)}_{=:\Phi} = \Phi^T \Phi,$$

where $\Phi = [\phi(\mathbf{x}_1), \phi(\mathbf{x}_2), \dots, \phi(\mathbf{x}_n)]$.



Mercer's Theorem

Let \mathcal{X} be a compact subset of \mathbb{R}^n and $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ a **kernel function**.

Then one can expand k in a uniformly convergent series of bounded functions ϕ_i s.t.

$$k(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \lambda_i \phi_i(\mathbf{x}) \phi_i(\mathbf{x}')$$

Mercer's Theorem

Let \mathcal{X} be a compact subset of \mathbb{R}^n and $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ a **kernel function**.

Then one can expand k in a uniformly convergent series of bounded functions ϕ_i s.t.

$$k(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \lambda_i \phi_i(\mathbf{x}) \phi_i(\mathbf{x}')$$

Can be generalized even further

Definition: kernel functions

- Data space \mathcal{X}
- A **kernel** is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ satisfying:

Definition: kernel functions

- Data space \mathcal{X}
- A **kernel** is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ satisfying:
 1. **Symmetry**: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{X}$ it must hold that
$$k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$$

Definition: kernel functions

- Data space \mathcal{X}
- A **kernel** is a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ satisfying:
 1. **Symmetry**: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{X}$ it must hold that $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$
 2. **Positive semi-definiteness**: For any n , any set $S = \{\mathbf{x}_1, \dots, \mathbf{x}_n\} \subseteq \mathcal{X}$, the kernel (Gram) matrix

$$\mathbf{K} = \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix}$$

must be positive semi-definite

Examples of kernels on \mathbb{R}^d

- Linear kernel:


$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$$

Examples of kernels on \mathbb{R}^d

- Linear kernel: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$
- Polynomial kernel: $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^d$


Examples of kernels on \mathbb{R}^d

- Linear kernel: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$
- Polynomial kernel: $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^d$
- Gaussian (RBF, squared exp. kernel): $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_2^2 / h^2)$




Examples of kernels on \mathbb{R}^d

- Linear kernel: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$
- Polynomial kernel: $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^d$
- Gaussian (RBF, squared exp. kernel): $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_2^2/h^2)$



- Laplacian kernel: $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_1/h)$



Examples of (non)-kernels

- $k(x, x') = \sin(x)\cos(x')$

Examples of (non)-kernels

- $k(x, x') = \sin(x)\cos(x')$

Not symmetric: Consider $x = 0, x' = \pi/2$:

$$k(x, x') = 0 \cdot 0 = 0 \neq 1 = 1 \cdot 1 = k(x', x)$$

Hence, k is not a valid kernel.

Examples of (non)-kernels

- $k(x, x') = \sin(x)\cos(x')$

Not symmetric: Consider $x = 0, x' = \pi/2$:

$$k(x, x') = 0 \cdot 0 = 0 \neq 1 = 1 \cdot 1 = k(x', x)$$

Hence, k is not a valid kernel.

- $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{M} \mathbf{x}'$

Examples of (non)-kernels

- $k(x, x') = \sin(x)\cos(x')$

Not symmetric: Consider $x = 0, x' = \pi/2$:

$$k(x, x') = 0 \cdot 0 = 0 \neq 1 = 1 \cdot 1 = k(x', x)$$

Hence, k is not a valid kernel.

- $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{M} \mathbf{x}'$

If \mathbf{M} is symmetric:

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{M} \mathbf{x}' = \mathbf{x}^T \mathbf{M}^T \mathbf{x}' = (\mathbf{x}^T \mathbf{M}^T \mathbf{x}')^T = \mathbf{x}'^T \mathbf{M} \mathbf{x} = k(\mathbf{x}', \mathbf{x})$$

If \mathbf{M} is not symmetric, then k not guaranteed to be symmetric.

Claim: k valid kernels $\Leftrightarrow \mathbf{M}$ psd

Effect of kernel on function class




- Given kernel k , predictors (for kernelized classification) have the form

$$\hat{y} = \text{sign} \left(\sum_{j=1}^n \alpha_j y_j k(\mathbf{x}_j, \mathbf{x}) \right)$$

Example: Gaussian kernel

$$k(\mathbf{x}, \mathbf{x}') = \exp \left(-\|\mathbf{x} - \mathbf{x}'\|_2^2 / h^2 \right)$$




$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$



Example: Laplace/Exponential kernel

$$k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_1/h)$$


$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$



Demo: Effect on decision boundary


- Can define kernels on a **variety of objects**:
 - Sequence kernels
 - Graph kernels
 - Diffusion kernels
 - Kernels on probability distributions
 - ...

Example: Graph kernels



- Can define a kernel for measuring similarity between graphs by comparing random walks on both graphs (not further defined here)

Example: Diffusion kernels on graphs



$$\mathbf{K} = \exp(-\beta \mathbf{L})$$

- Can measure similarity among nodes in a graph via diffusion kernels (not defined here)

Kernel engineering (composition rules)

- Suppose we have two kernels:

$$k_1: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R} \quad k_2: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

defined on data space \mathcal{X}

Kernel engineering (composition rules)

- Suppose we have two kernels:

$$k_1: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R} \quad k_2: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

defined on data space \mathcal{X}

- Then the following functions are valid kernels:

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = ck_1(\mathbf{x}, \mathbf{x}') \text{ for } c > 0$$

$$k(\mathbf{x}, \mathbf{x}') = f(k_1(\mathbf{x}, \mathbf{x}')),$$

where f is a polynomial with positive coefficients or the exponential function

Example: ANOVA kernel

$$\text{General: } k_P(\mathbf{x}, \mathbf{x}') = \sum_{1 \leq j_1 < \dots < j_P \leq d} \prod_{p=1}^P k_{j_p}(x_{j_p}, x'_{j_p}) \quad \mathbf{x}, \mathbf{x}' \in \mathbb{R}^d$$



$$\text{P=1: } \underbrace{k(\mathbf{x}, \mathbf{x}')}_{k: \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}} = \sum_{j=1}^d \underbrace{k_j(x_j, x'_j)}_{k_j: \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}} \quad \mathbf{x}, \mathbf{x}' \in \mathbb{R}^d$$

The function k is a valid kernel according to composition rules. Which functions f are learned?

$$\begin{aligned} f(\mathbf{x}) &= \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}^{(i)}, \mathbf{x}) \\ &= \sum_{i=1}^n \alpha_i y_i \sum_{j=1}^d k_j(x_j^{(i)}, x_j) = \sum_{j=1}^d \underbrace{\sum_{i=1}^n \alpha_i y_i k_j(x_j^{(i)}, x_j)}_{f_j(x_j)} \end{aligned}$$

Example: Modeling pairwise data

- May want to use kernels to model pairwise data (users \times products; genes \times patients; ...)



$$k((x, z), (x', z')) = k_{\text{user}}(x, x') \cdot k_{\text{product}}(z, z')$$

$$k((x, z), (x', z')) = k_{\text{user}}(x, x') + k_{\text{product}}(z, z')$$

Where are we?

- We've seen how to kernelize the perceptron
- Discussed properties of kernels, and seen examples

- We've seen how to kernelize the perceptron
- Discussed properties of kernels, and seen examples
- Next questions:
 - What kind of predictors / decision boundaries do kernel methods entail?
 - Can we use the kernel trick beyond the perceptron?

Kernels as similarity functions

- Recall Perceptron (and SVM) classification rule:


$$y = \text{sign} \left(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) \right)$$

Kernels as similarity functions

- Recall Perceptron (and SVM) classification rule:

$$y = \text{sign} \left(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) \right)$$

- Consider Gaussian kernel $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2/h^2)$




Kernels as similarity functions

- Recall Perceptron (and SVM) classification rule:

$$y = \text{sign} \left(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) \right) \approx \text{sign} \left(\sum_{i=1}^n \alpha_i y_i [\mathbf{x}_i \text{ "close" to } \mathbf{x}] \right)$$

- Consider Gaussian kernel $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2/h^2)$



Side note: Nearest-neighbor classifiers

- For data point \mathbf{x} , predict majority of labels of k nearest neighbors:

$$y = \text{sign} \left(\sum_{i=1}^n y_i [\mathbf{x}_i \text{ among } k \text{ nearest neighbors } \mathbf{x}] \right)$$

Nearest-neighbor classifiers

- For data point \mathbf{x} , predict majority of labels of k nearest neighbors:

$$y = \text{sign} \left(\sum_{i=1}^n y_i [\mathbf{x}_i \text{ among } k \text{ nearest neighbors } \mathbf{x}] \right)$$

❓ How to choose k ?

Nearest-neighbor classifiers

- For data point \mathbf{x} , predict majority of labels of k nearest neighbors:

$$y = \text{sign} \left(\sum_{i=1}^n y_i [\mathbf{x}_i \text{ among } k \text{ nearest neighbors } \mathbf{x}] \right)$$

❓ How to choose k ? 💡 Cross-validation!

- k -nearest neighbor:

$$y = \text{sign} \left(\sum_{i=1}^n y_i [\mathbf{x}_i \text{ among } k \text{ nearest neighbors } \mathbf{x}] \right)$$

- Kernel perceptron:

$$y = \text{sign} \left(\sum_{i=1}^n y_i \alpha_i k(\mathbf{x}_i, \mathbf{x}) \right)$$

Comparison: k-NN vs Kernelized Perceptron

Method	<i>k-NN</i>	<i>Kernelized Perceptron</i>
Advantages	No training necessary	Optimized weights can lead to improved performance Can capture “global trends” with suitable kernels Depends on “wrongly classified” examples only
Disadvantages	Depends on all data ⇒ inefficient	Training requires optimization

Parametric vs nonparametric learning

- **Parametric models** have finite set of parameters
- **Example:** Linear regression, linear Perceptron, . . .

Parametric vs nonparametric learning

- **Parametric models** have finite set of parameters
- **Example:** Linear regression, linear Perceptron, ...
- **Nonparametric models** grow in complexity with the size of the data
 - Potentially much more expressive
 - But also more computationally complex. Why?
- **Example:** Kernelized Perceptron, k-NN, ...

Parametric vs nonparametric learning

- **Parametric models** have finite set of parameters
- **Example:** Linear regression, linear Perceptron, ...
- **Nonparametric models** grow in complexity with the size of the data
 - Potentially much more expressive
 - But also more computationally complex. Why?
- **Example:** Kernelized Perceptron, k-NN, ...
- 💡 **Kernels provide a principled way of deriving nonparametric models from parametric ones**

Where are we?

- We've seen how to kernelize the perceptron
- Discussed properties of kernels, and seen examples

Where are we?

- We've seen how to kernelize the perceptron
- Discussed properties of kernels, and seen examples
- Next questions:
 - Can we use the kernel trick beyond the perceptron?

- The support vector machine

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \mathbf{w}^T \mathbf{x}_i\} + \lambda \|\mathbf{w}\|_2^2$$

can also be kernelized

How to kernelize the objective?

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \mathbf{w}^T \mathbf{x}_i\} + \lambda \|\mathbf{w}\|_2^2$$

$$\begin{aligned} \max\{0, 1 - y_i \mathbf{w}^T \mathbf{x}_i\} &= \max\left\{0, 1 - y_i \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j\right)^T \mathbf{x}_i\right\} \\ &= \max\left\{0, 1 - y_i \sum_{j=1}^n \alpha_j y_j \underbrace{\mathbf{x}_j^T \mathbf{x}_i}_{=k(\mathbf{x}_j, \mathbf{x}_i)}\right\} \\ &= \max\left\{0, 1 - y_i \alpha^T \mathbf{k}^{(i)}\right\}, \end{aligned}$$

where $\alpha^T = (\alpha_1, \dots, \alpha_n)$, $\mathbf{k}^{(i)} = (y_1 k(\mathbf{x}_1, \mathbf{x}_i), \dots, y_n k(\mathbf{x}_n, \mathbf{x}_i))^T$.

How to kernelize the regularizer?

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \mathbf{w}^T \mathbf{x}_i\} + \lambda \|\mathbf{w}\|_2^2$$

$$\begin{aligned} \lambda \|\mathbf{w}\|_2^2 &= \lambda \mathbf{w}^T \mathbf{w} = \lambda \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right)^T \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j \right) \\ &= \lambda \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \underbrace{\mathbf{x}_i^T \mathbf{x}_j}_{=k(\mathbf{x}_i, \mathbf{x}_j)} = \lambda \alpha^T \mathbf{D}_y \mathbf{K} \mathbf{D}_y \alpha, \end{aligned}$$

$$\text{where } \mathbf{D}_y = \begin{pmatrix} y_1 & & \\ & \ddots & \\ & & y_n \end{pmatrix} \text{ and } \mathbf{K} = \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_1, \mathbf{x}_n) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix}$$

Learning

Solve the following problem.

- **Perceptron:** $\operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -y_i \alpha^T \mathbf{k}_i\}$
- **SVM:** $\operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \alpha^T \mathbf{k}_i\} + \lambda \alpha^T \mathbf{D}_y \mathbf{K} \mathbf{D}_y \alpha$

$$\mathbf{k}_i = [y_1 k(\mathbf{x}_i, \mathbf{x}_1), \dots, y_n k(\mathbf{x}_i, \mathbf{x}_n)]$$

Learning & prediction with kernel classifier

Learning

Solve the following problem.

- **Perceptron:** $\operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, -y_i \alpha^T \mathbf{k}_i\}$
- **SVM:** $\operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \alpha^T \mathbf{k}_i\} + \lambda \alpha^T \mathbf{D}_y \mathbf{K} \mathbf{D}_y \alpha$

$$\mathbf{k}_i = [y_1 k(\mathbf{x}_i, \mathbf{x}_1), \dots, y_n k(\mathbf{x}_i, \mathbf{x}_n)]$$

Prediction



For data point \mathbf{x} predict label y as

$$\hat{y} = \operatorname{sign} \left(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) \right)$$

Demo: Kernelized SVM

Kernelized Linear Regression

- From linear to **nonlinear** regression:



- Can also kernelize linear regression
- Predictor has the form

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x})$$

Example: Kernelized linear regression

- Original (**parametric**) linear optimization problem:

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2$$

- Similar as in perceptron, optimal $\hat{\mathbf{w}}$ lies in span of data:

$$\hat{\mathbf{w}} = \sum_{i=1}^n \alpha_i \mathbf{x}_i$$

Kernelizing linear regression

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2$$

$$\mathbf{w}^T \mathbf{x}_i - y_i = \sum_{j=1}^n \alpha_j \mathbf{x}_j^T \mathbf{x}_i - y_i = \sum_{j=1}^n \alpha_j k(\mathbf{x}_j, \mathbf{x}_i) - y_i = \alpha^T \mathbf{k}_i - y_i$$

$$\lambda \|\mathbf{w}\|_2^2 = \lambda \alpha^T \mathbf{K} \alpha,$$

where $\mathbf{k}_i = (k(\mathbf{x}_1, \mathbf{x}_i), \dots, k(\mathbf{x}_n, \mathbf{x}_i))^T$ and $\mathbf{K} = [\mathbf{k}_1 | \dots | \mathbf{k}_n]$.

$$\hat{\alpha} = \underset{\alpha}{\operatorname{argmin}} \frac{1}{n} \underbrace{\sum_{i=1}^n (\alpha^T \mathbf{k}_i - y_i)^2}_{= \|\alpha^T \mathbf{K} - \mathbf{y}\|_2^2 \text{ with } \mathbf{y} = (y_1, \dots, y_n)^T} + \lambda \alpha^T \mathbf{K} \alpha$$

- **Learning:** Solve least squares problem

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \|\alpha^T \mathbf{K} - \mathbf{y}\|_2^2 + \lambda \alpha^T \mathbf{K} \alpha$$

- **Learning:** Solve least squares problem

$$\hat{\alpha} = \operatorname{argmin}_{\alpha} \frac{1}{n} \|\alpha^T \mathbf{K} - \mathbf{y}\|_2^2 + \lambda \alpha^T \mathbf{K} \alpha$$

Closed-form solution: $\hat{\alpha} = (\mathbf{K} + n\lambda \mathbf{I})^{-1} \mathbf{y}$

- **Prediction:** For data point \mathbf{x} predict response y as

$$\hat{y} = \sum_{i=1}^n \hat{\alpha}_i k(\mathbf{x}_i, \mathbf{x})$$

Demo: Kernelized linear regression

KLR for the linear kernel

- What if $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$?

KLR for the linear kernel

- What if $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$?

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}) = \left(\sum_{i=1}^n \alpha_i \mathbf{x}_i^T \right) \mathbf{x} = \mathbf{w}^T \mathbf{x}$$

Application: Semi-parametric regression

- Often, parametric models are too “rigid”, and nonparametric models fail to extrapolate
- **Solution:** Use additive combination of linear and nonlinear kernel function:


$$k(\mathbf{x}, \mathbf{x}') = c_1 \exp(-\|\mathbf{x} - \mathbf{x}'\|_2^2/h^2) + c_2 \mathbf{x}^T \mathbf{x}'$$

$$\begin{aligned} f(\mathbf{x}) &= \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}) = \sum_{i=1}^n \alpha_i (c_1 \exp(-\|\mathbf{x}_i - \mathbf{x}\|_2^2/h^2) + c_2 \mathbf{x}_i^T \mathbf{x}) \\ &= c_1 \sum_{i=1}^n \alpha_i \exp(-\|\mathbf{x}_i - \mathbf{x}\|_2^2/h^2) + c_2 \sum_{i=1}^n \alpha_i \mathbf{x}_i^T \mathbf{x} \\ &= c_1 f_1(\mathbf{x}) + c_2 f_2(\mathbf{x}) \end{aligned}$$


Demo: Semi-parametric KLR

Example

Function




Training data




Example fits


Linear kernel




Gaussian kernel



Gaussian + linear kernel





Periodic + linear kernel




- For a given kernel, how should we choose parameters?

Choosing kernels

- For a given kernel, how should we choose parameters?
 -  Cross-validation

- For a given kernel, how should we choose parameters?
 -  Cross-validation
- How should we select suitable kernels?
 - Domain knowledge (dependent on data type)
 - “Brute force” (or heuristic) search
 - Use cross-validation

- For a given kernel, how should we choose parameters?
 -  Cross-validation
- How should we select suitable kernels?
 - Domain knowledge (dependent on data type)
 - “Brute force” (or heuristic) search
 - Use cross-validation
- Learning kernels
 - Much research on automatically selecting good kernels (Multiple Kernel Learning; Hyperkernels; etc.)

What about overfitting?

- Kernels map to (very) high-dimensional spaces.
- Why do we hope to be able to learn?
- **First attempt of an answer:**
(typically) # parameters \ll # dimensions. Why?

What about overfitting?

- Kernels map to (very) high-dimensional spaces.
- Why do we hope to be able to learn?
- **First attempt of an answer:**
(typically) # parameters \ll # dimensions. Why?
- Number of parameters = number of data points
("non-parametric learning")

What about overfitting?

- Kernels map to (very) high-dimensional spaces.
- Why do we hope to be able to learn?
- **Second attempt of an answer:**

What about overfitting?

- Kernels map to (very) high-dimensional spaces.
- Why do we hope to be able to learn?
- **Second attempt of an answer:**
- Overfitting can of course happen (if we choose poor parameters)

What about overfitting?

- Kernels map to (very) high-dimensional spaces.
- Why do we hope to be able to learn?
- **Second attempt of an answer:**
- Overfitting can of course happen (if we choose poor parameters)
- Can combat overfitting by regularization:
 - This is already built into kernelized linear regression (and SVMs), but not the kernelized Perceptron

$$\text{KLR: } \operatorname{argmin}_{\alpha} \frac{1}{n} \|\alpha^T \mathbf{K} - \mathbf{y}\|_2^2 + \lambda \alpha^T \mathbf{K} \alpha$$

$$\text{SVM: } \operatorname{argmin}_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \alpha^T \mathbf{k}_i\} + \lambda \alpha^T \mathbf{D}_y \mathbf{K} \mathbf{D}_y \alpha$$

Before we conclude

- Our proof for why we can use the ansatz $\hat{\mathbf{w}} = \sum_i \alpha_i \phi(\mathbf{x}_i)$ was hand-wavy
- Principled proofs: **Representer Theorems**

Before we conclude

- Our proof for why we can use the ansatz $\hat{\mathbf{w}} = \sum_i \alpha_i \phi(\mathbf{x}_i)$ was hand-wavy
- Principled proofs: **Representer Theorems**

Representer Theorem

Let \mathcal{X} be the data space and $\phi: \mathcal{X} \rightarrow \mathcal{H}$ a mapping from \mathcal{X} to Hilbert space \mathcal{H} . Consider the optimization problem

$$\min_{\mathbf{w}} f(\langle \mathbf{w}, \phi(\mathbf{x}_1) \rangle, \dots, \langle \mathbf{w}, \phi(\mathbf{x}_n) \rangle) + R(\|\mathbf{w}\|),$$

where f is an arbitrary function from $f: \mathbb{R}^n \rightarrow \mathbb{R}$ and $R: \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}$ is a strictly monotonic function. Then, there exist $\alpha_1, \dots, \alpha_n$ such that $\mathbf{w} = \sum_{i=1}^n \alpha_i \phi(\mathbf{x}_i)$ is an optimal solution for the above optimization problem.

Representer Theorem I

Many functions satisfy the above form:

- Ridge regression with non-linear functions:

$$f(\underbrace{\langle \mathbf{w}, \phi(\mathbf{x}_1) \rangle}_{\hat{y}_1}, \dots, \underbrace{\langle \mathbf{w}, \phi(\mathbf{x}_n) \rangle}_{\hat{y}_n}) = \sum_{i=1}^n (\hat{y}_i - y_i)^2 \quad R(\|\mathbf{w}\|) = \|\mathbf{w}\|_2^2$$

Proof of the Representer Theorem:

Let \mathbf{w}^* be an optimal solution of the problem. Since $\mathbf{w}^* \in \mathcal{H}$,

$$\mathbf{w}^* = \underbrace{\sum_{i=1}^m \alpha_i \phi(\mathbf{x}_i)}_{\text{subspace of } \mathcal{H}} + \mathbf{u},$$

where \mathbf{u} is orthogonal to the subspace, i.e., $\langle \mathbf{u}, \phi(\mathbf{x}_i) \rangle = 0$ for all $i \in \{1, \dots, m\}$.

Representer Theorem II

We prove the theorem by showing that the same value of the objective function can be reached if $\mathbf{u} = \mathbf{0}$.

Let $\mathbf{w} = \mathbf{w}^* - \mathbf{u}$.

Regularizer:

$$\|\mathbf{w}^*\|^2 = \|\mathbf{w}\|^2 + \|\mathbf{u}\|^2 \Rightarrow \|\mathbf{w}\| \leq \|\mathbf{w}^*\| \Rightarrow R(\|\mathbf{w}\|) \leq R(\|\mathbf{w}^*\|)$$

f -term:

$$\begin{aligned}\langle \mathbf{w}, \phi(\mathbf{x}_i) \rangle &= \langle \mathbf{w}^* - \mathbf{u}, \phi(\mathbf{x}_i) \rangle = \langle \mathbf{w}^*, \phi(\mathbf{x}_i) \rangle, \text{ since } \langle \mathbf{u}, \phi(\mathbf{x}_i) \rangle = 0 \\ f(\langle \mathbf{w}, \phi(\mathbf{x}_1) \rangle, \dots, \langle \mathbf{w}, \phi(\mathbf{x}_n) \rangle) &= f(\langle \mathbf{w}^*, \phi(\mathbf{x}_1) \rangle, \dots, \langle \mathbf{w}^*, \phi(\mathbf{x}_n) \rangle)\end{aligned}$$

Representer Theorem III

We have shown that the value of the objective function for \mathbf{w} is at most the value for \mathbf{w}^* . Hence,

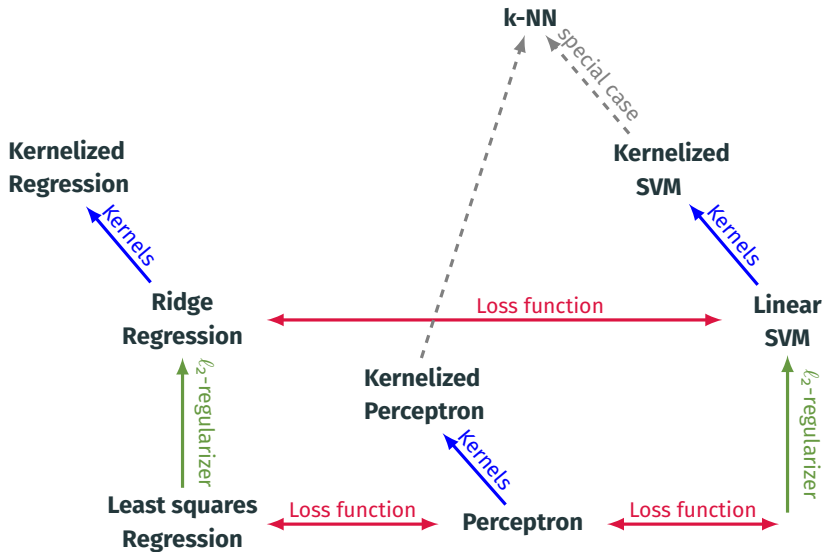
$$\mathbf{w} = \sum_{i=1}^n \alpha_i \phi(\mathbf{x}_i)$$

also is an optimal solution, proving the theorem. □

What you need to know

- Kernels are
 - (efficient, implicit) inner products
 - Positive (semi-)definite functions
 - Many examples (linear, polynomial, Gaussian/RBF, ...)
- The “Kernel trick”
 - Reformulate learning algorithm so that inner products appear
 - Replace inner products by kernels
- K-Nearest Neighbor classifier (and relation to Perceptron)
- How to choose kernels (kernel engineering etc.)
- Applications: Kernelized Perceptron / SVM; kernelized linear regression

Supervised learning big picture so far



Supervised learning summary so far

Representation/
features

Linear hypotheses, non-linear hypotheses through feature transformations, kernels

Model/
objective

Loss-function (squared loss, ℓ_p loss, 0/1 loss, Perceptron loss, Hinge loss) + Regularization (ℓ_2 norm)

Method

Exact solution, Gradient Descent, (mini-batch) SGD

Evaluation
metric

Empirical risk = (mean) squared error, Accuracy

Model
selection

k -fold cross-validation, Monte Carlo cross-validation

- S. Shalev-Schwartz & S. Ben-David, “Understanding Machine Learning: From Theory to Algorithms”, Chapter 16
- C. Bishop, “Pattern Recognition and Machine Learning”, Chapter 6 (6.1 & 6.2)